



Benchmarking Monocular Depth Estimation Models

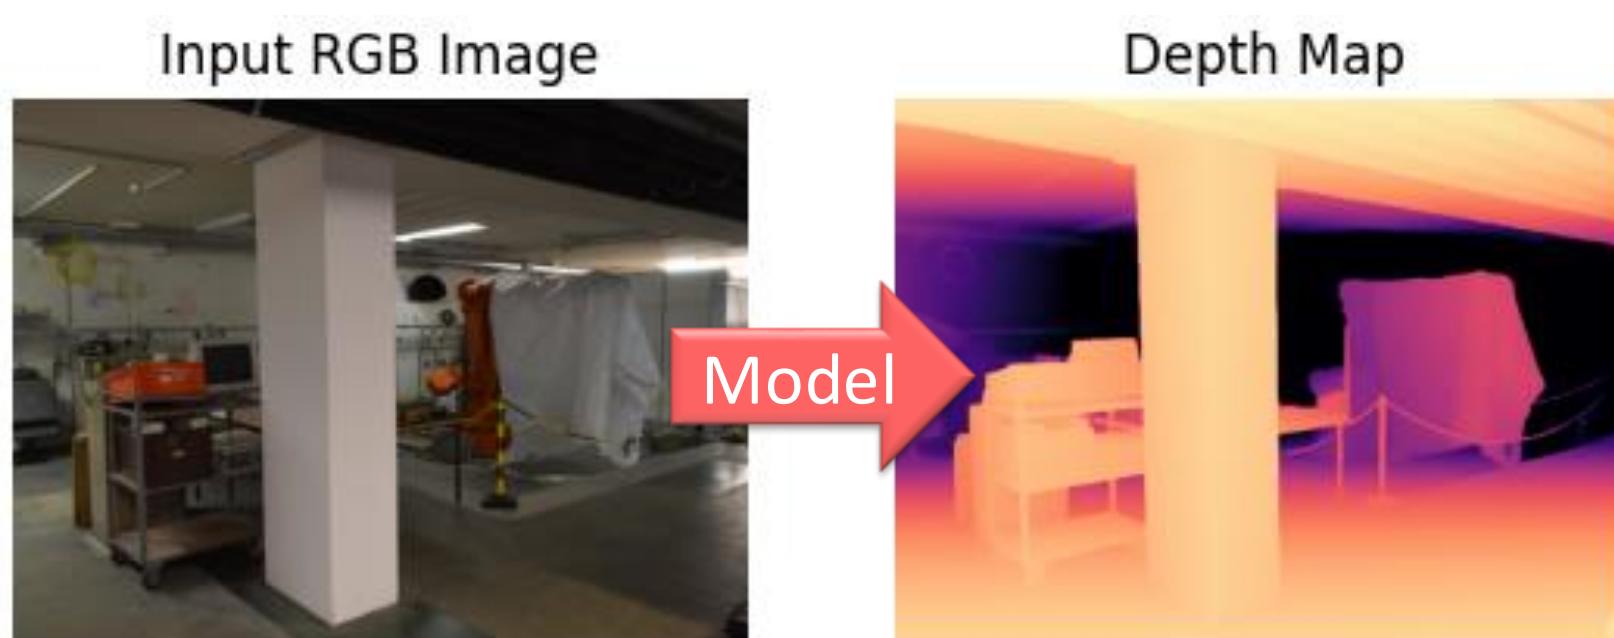


Nikolas Spyropoulos, el21202, el21202@mail.ntua.gr

School of Electrical and Computer Engineering, National Technical University of Athens

Introduction & Motivation

- Monocular Depth Estimation (MDE) is critical for **3D Reconstruction, AR, and Autonomous Driving**
- The Problem:** Models struggle with **Domain Shift**.
Models trained on indoor scenes fail to generalize to outdoor environments
- The Goal:** Assess the **Zero-Shot Domain Generalization** of modern Foundation Models across diverse environments

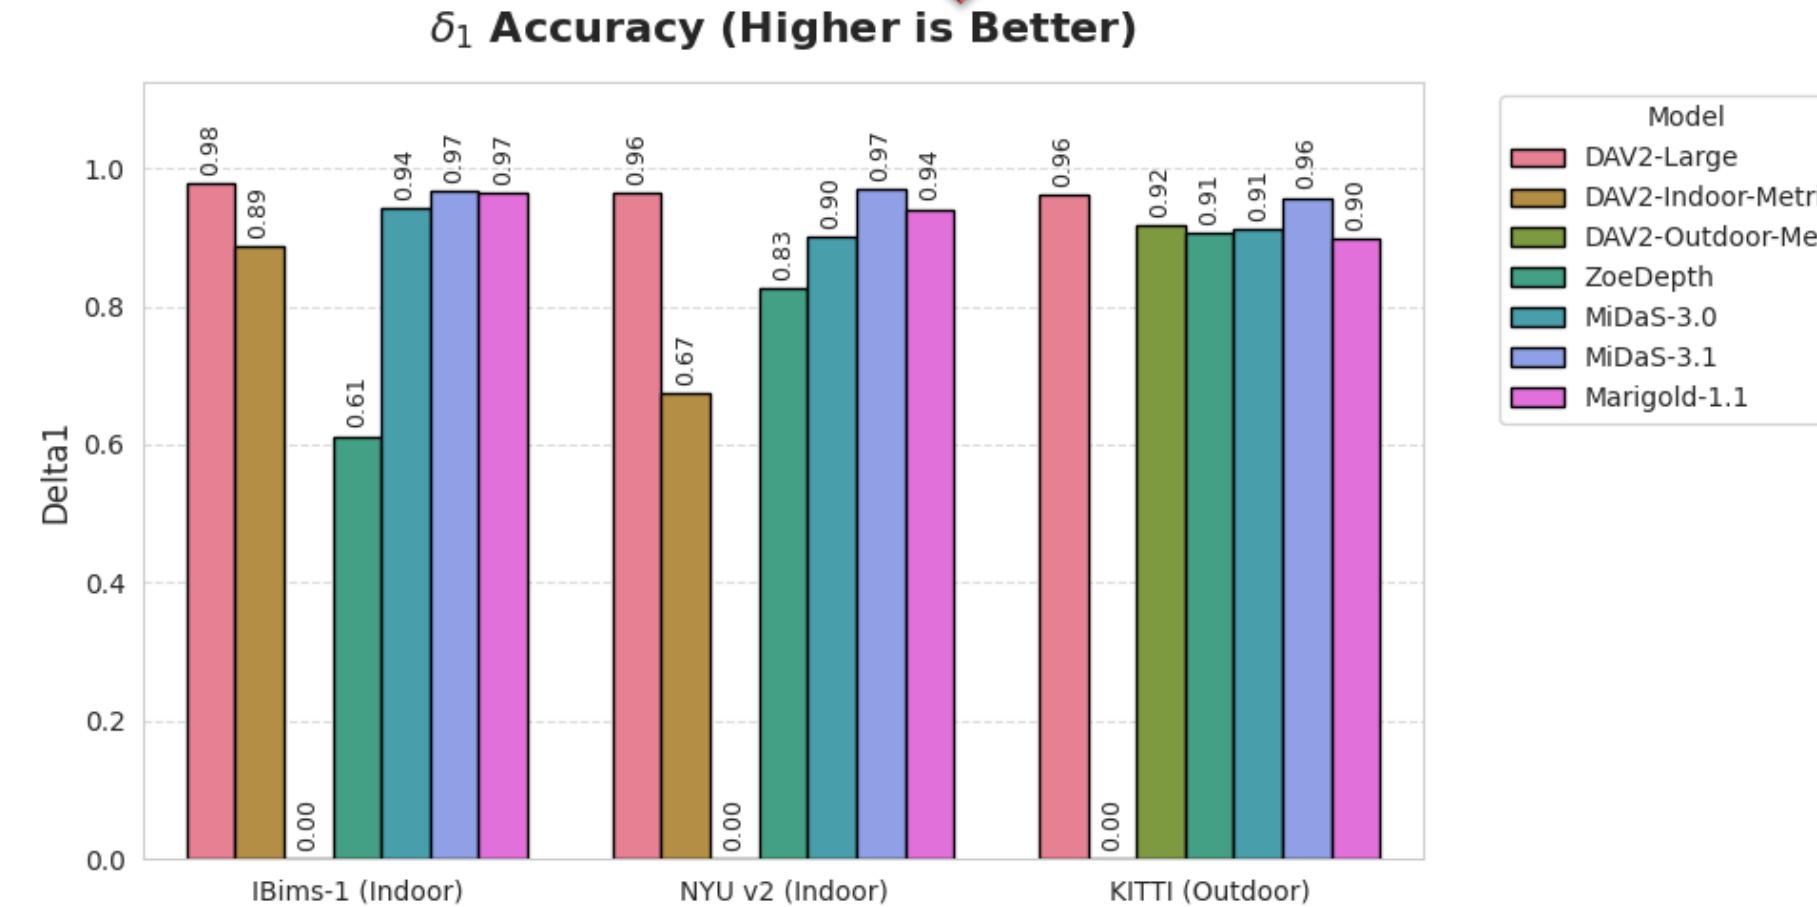
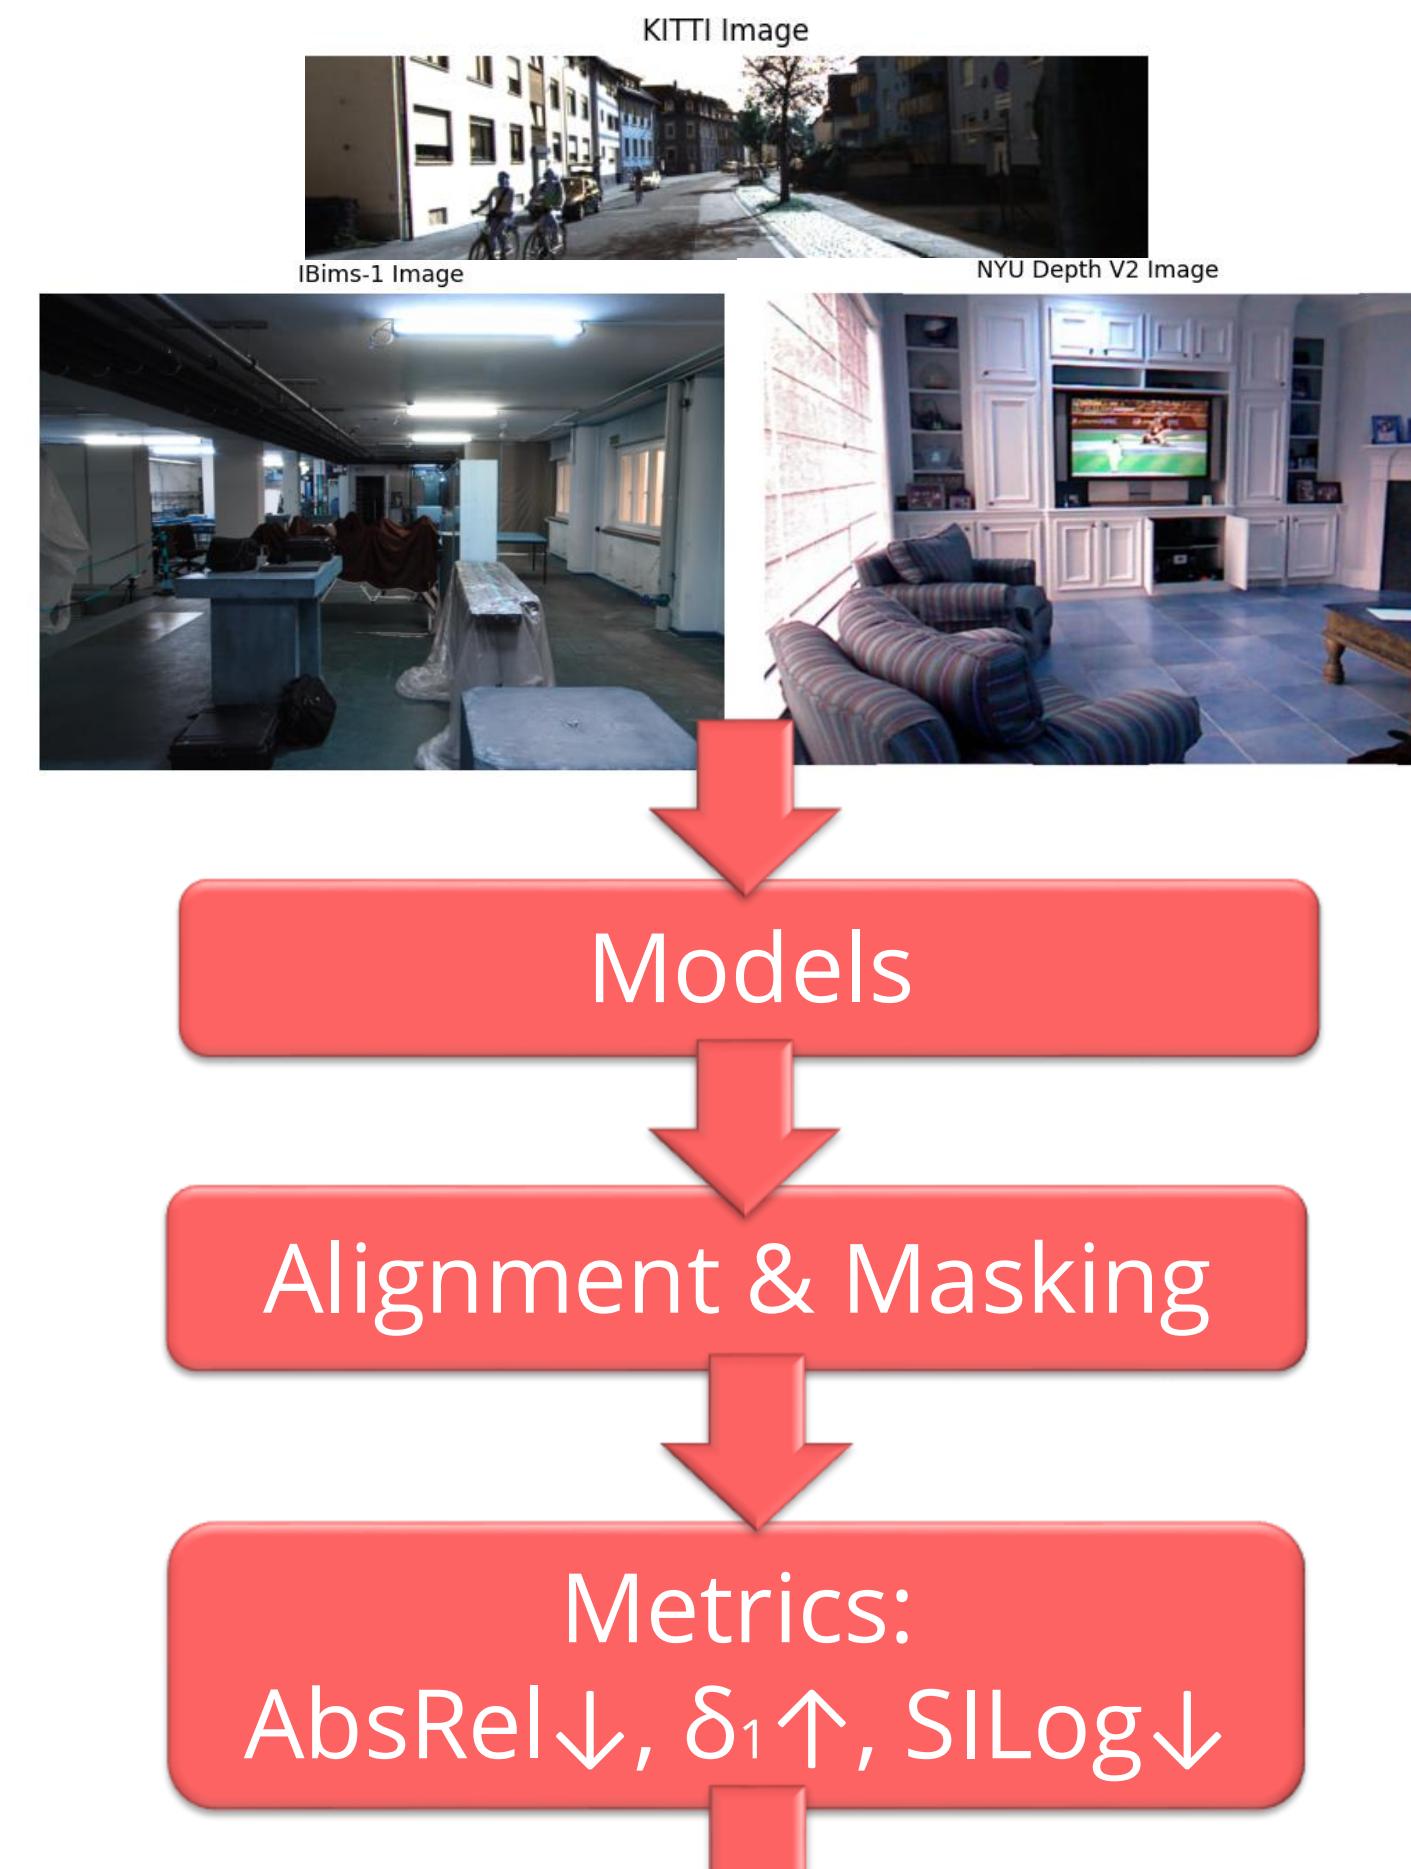


The Models

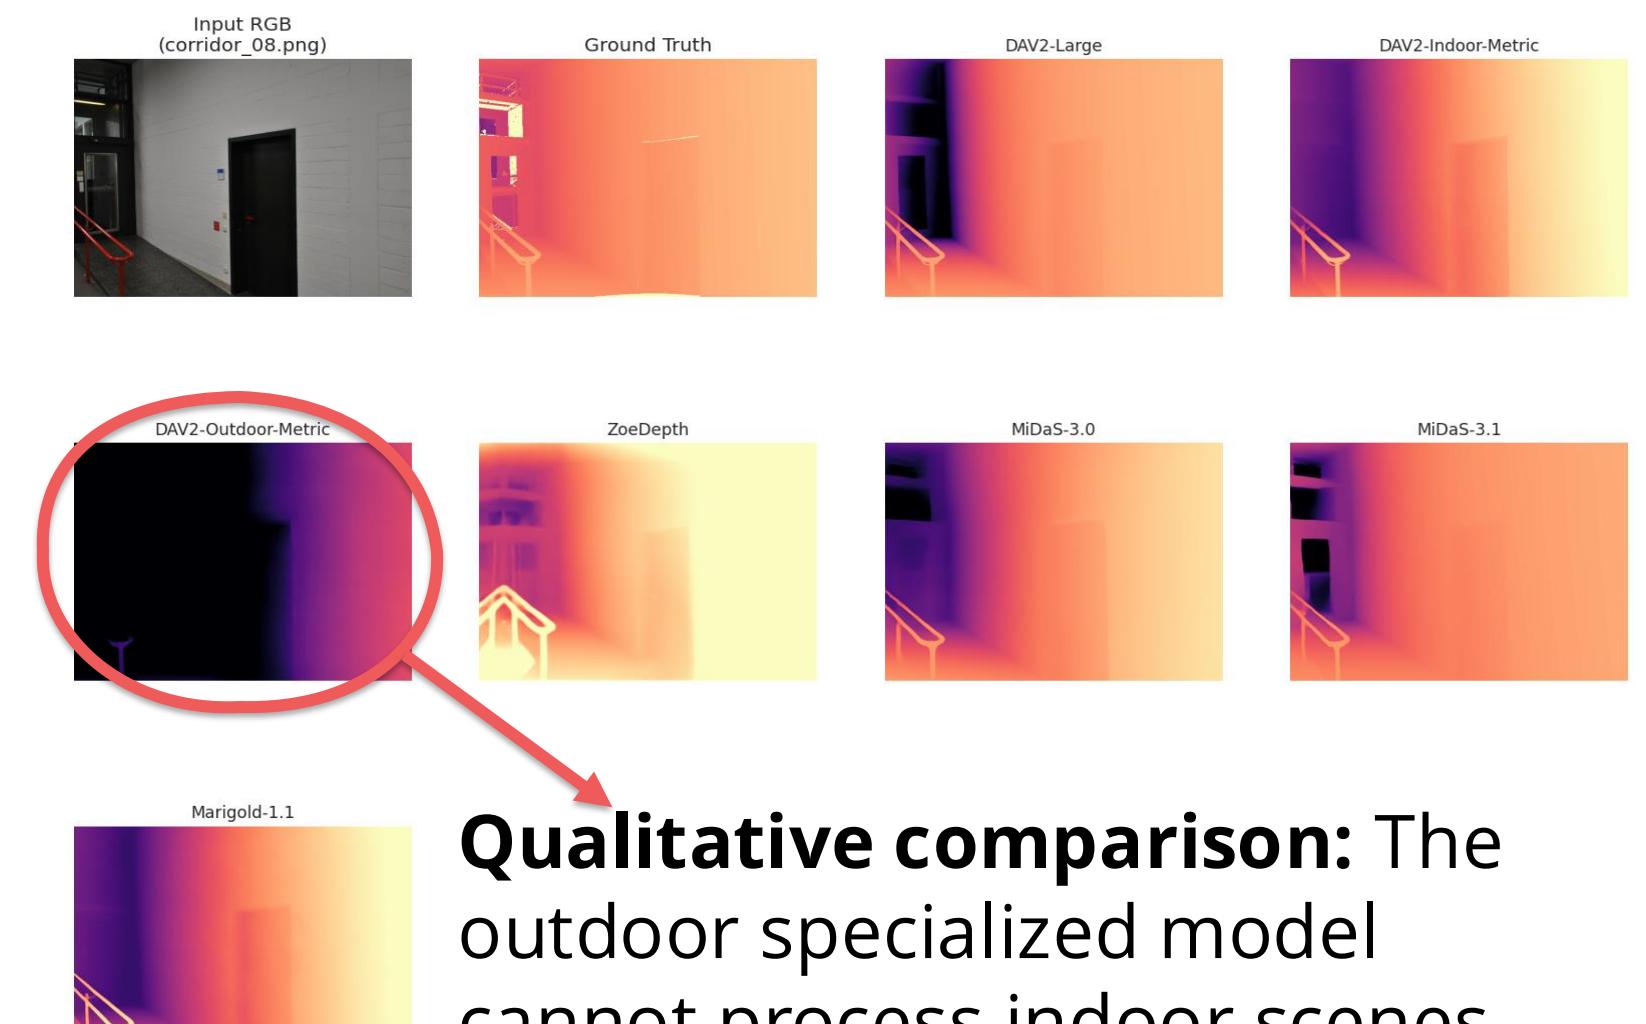
- Depth Anything V2:** Relative Depth, Transformer
 - DAV2-Large:** basic large model
 - DAV2-Indoor-Metric:** fine-tuned for indoor scenes and metric depth
 - DAV2-Outdoor-Metric:** fine-tuned for outdoor scenes and metric depth
- ZoeDepth:** Metric Depth, Transformer
- MiDaS v3.0 & v3.1:** Relative Depth, Transformer
- Marigold v1.1:** Relative Depth, Stable Diffusion

The Pipeline

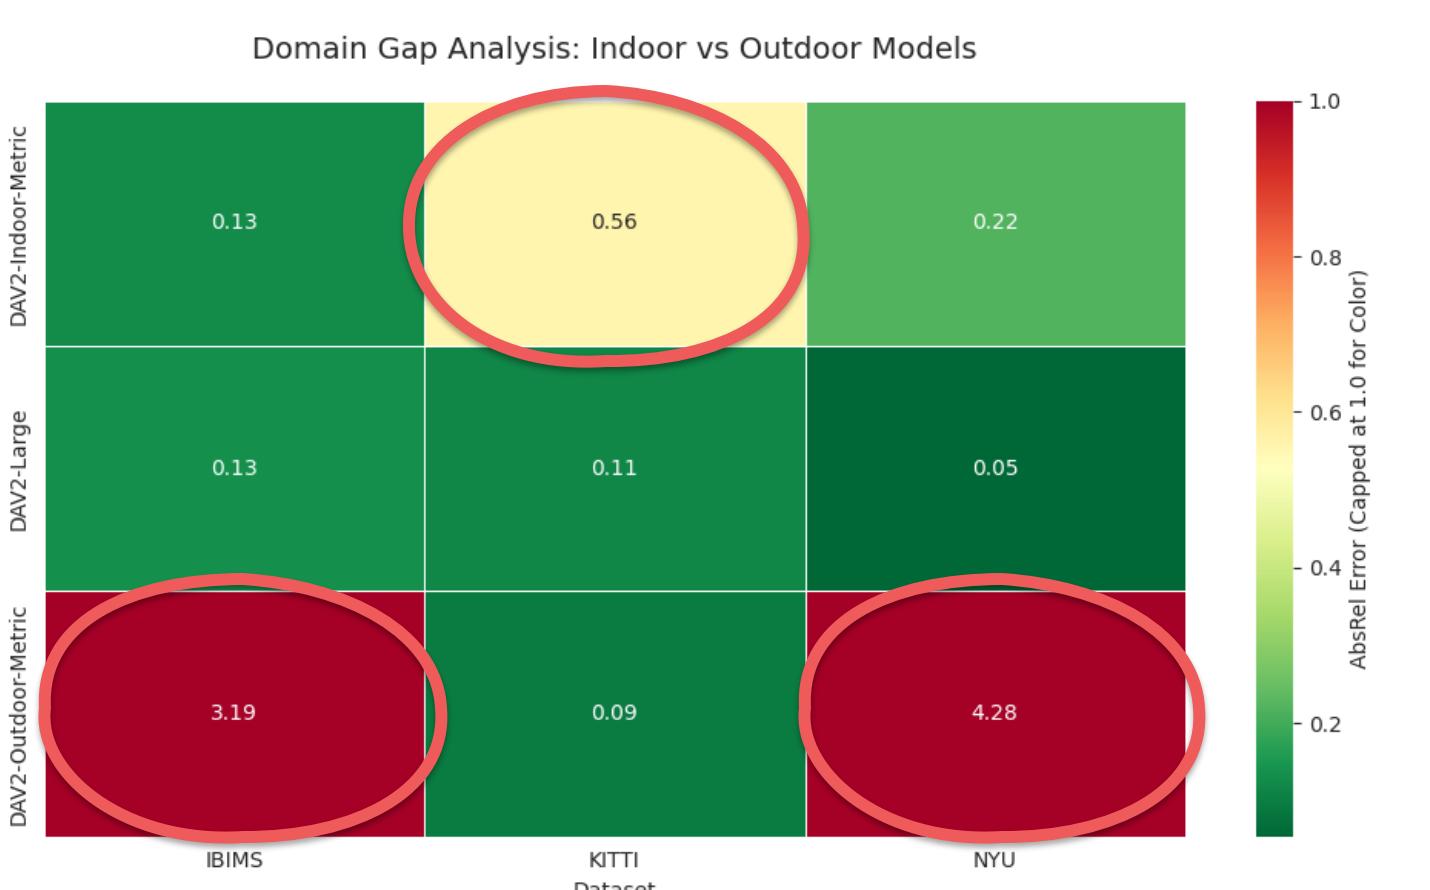
Datasets: IBims-1: Indoor/Challenging, NYU: Indoor, KITTI: Outdoor



Domain Gap



Qualitative comparison: The outdoor specialized model cannot process indoor scenes



Conclusions

- DAV2** is the **best all-rounder** model with **real-time** inference time
- Marigold** captures the **finest details**, but it cannot be used in real-time
- Relative depth models** are superior for "**in-the-wild**" scenarios
- Specialized** models **lack domain generalization** abilities